

8E and 8F: Finding the Probability $P(Y=1|X)$

8E: Implementing Decision Function of SVM RBF Kernel

After we train a kernel SVM model, we will be getting support vectors and their corresponding coefficients a_i . Check the documentation for better understanding of these attributes:

<https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html>

```
Attributes: support_ : array-like, shape = [n_SV]
            Indices of support vectors.

            support_vectors_ : array-like, shape = [n_SV, n_features]
            Support vectors.

            n_support_ : array-like, dtype=int32, shape = [n_class]
            Number of support vectors for each class.

            dual_coef_ : array, shape = [n_class-1, n_SV]
            Coefficients of the support vector in the decision function. For multiclass, coefficient for all 1-vs-1
            classifiers. The layout of the coefficients in the multiclass case is somewhat non-trivial. See the
            section about multi-class classification in the SVM section of the User Guide for details.

            coef_ : array, shape = [n_class * (n_class-1) / 2, n_features]
            Weights assigned to the features (coefficients in the primal problem). This is only available in the
            case of a linear kernel.

            coef_ is a readonly property derived from dual_coef_ and support_vectors_.

            intercept_ : array, shape = [n_class * (n_class-1) / 2]
            Constants in decision function.

            fit_status_ : int
            0 if correctly fitted, 1 otherwise (will raise warning)

            probA_ : array, shape = [n_class * (n_class-1) / 2]
            probB_ : array, shape = [n_class * (n_class-1) / 2]
            If probability=True, the parameters learned in Platt scaling to produce probability estimates from
            decision values. If probability=False, an empty array. Platt scaling uses the logistic function
             $1 / (1 + \exp(\text{decision\_value} * \text{probA\_} + \text{probB\_}))$  where probA_ and probB_ are learned
            from the dataset [R20c70293ef72-2]. For more information on the multiclass case and training
            procedure see section 8 of [R20c70293ef72-1].
```

As a part of this assignment you will be implementing the `decision_function()` of kernel SVM, here `decision_function()` means based on the value return by `decision_function()` model will classify the data point either as positive or negative

Ex 1: In logistic regression After training the models with the optimal weights w we get, we will find the value $1/(1+\exp(-(wx+b)))$, if this value comes out to be < 0.5 we will mark it as negative class, else its positive class

Ex 2: In Linear SVM After training the models with the optimal weights w we get, we will find the value of $\text{sign}(wx + b)$, if this value comes out to be -ve we will mark it as negative class, else its positive class.

Similarly in Kernel SVM After training the models with the coefficients a_i we get, we will find the value of $\text{sign}(\sum_{i=1} (y_i a_i K(x_i, x_q)) + \text{intercept})$, here $K(x_i, x_q)$ is the RBF kernel. If this value comes out to be -ve we will mark x_q as negative class, else its positive class.

RBF kernel is defined as: $K(x_i, x_q) = \exp(-\gamma ||x_i - x_q||^2)$

For better understanding check this link: <https://scikit-learn.org/stable/modules/svm.html#svm-mathematical-formulation>

Task E

1. Split the data into $X_{\text{train}}(60)$, $X_{\text{cv}}(20)$, $X_{\text{test}}(20)$
2. Train SVC($\gamma = 0.001, C = 100$.) on the $(X_{\text{train}}, Y_{\text{train}})$
3. Get the decision boundary values f_{cv} on the X_{cv} data i.e. `f_cv = decision_function(X_cv)`

you need to implement this decision_function()

In [82]:

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.svm import SVC
```

In [83]:

```
X, y = make_classification(n_samples=5000, n_features=5, n_redundant=2,
                          n_classes=2, weights=[0.7], class_sep=0.7, random_state=1
5)
X.shape
```

Out[83]:

```
(5000, 5)
```

Pseudo code

```
clf = SVC(gamma=0.001, C=100.)
clf.fit(Xtrain, ytrain)
```

```
def decision_function(Xcv, ...): #use appropriate parameters
    for a data point  $x_q$  in Xcv:
        #write code to implement  $(\sum_{\text{all the support vectors } i=1} (y_i a_i K(x_i, x_q)) + \text{intercept})$ , here the values  $y_i, a_i$ ,
and intercept can be obtained from the trained model
    return # the decision_function output for all the data points in the Xcv
```

```
fcv = decision_function(Xcv, ...) # based on your requirement you can pass any other parameters
```

Note: Make sure the values you get as fcv, should be equal to outputs of `clf.decision_function(Xcv)`

In [84]:

```
# you can write your code here
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.2)

X_tr, X_cv, Y_tr, Y_cv = train_test_split(X_train, Y_train, test_size=0.25)
```

In [85]:

```
def decision_function(Xcv, sv, y_alp, intercept, gam):
    dec_fn = []
    #print(sv)
    #print(gam)
    #print(len(Xcv))
    for i in range(len(Xcv)):
        sum = 0
        for j in range(len(sv)):
            sum += (y_alp[j]*np.exp(-1*np.sum((Xcv[i,:]-sv[j,:])**2)*gam))
        dec_fn.append((sum+intercept).tolist())

    dec_fn = [i[0] for i in dec_fn]
    return dec_fn
```

In [86]:

```
gam = 0.001
clf = SVC(gamma= gam, C=100)
clf.fit(X_tr, Y_tr)
```

```

print(len(clf.support_))
y_alpha = clf.dual_coef_
intercept = clf.intercept_
sv = clf.support_vectors_
#print((y_alpha))
dec_fn_imlem = decision_function(X_cv,sv,y_alpha[0],intercept,gam)
dec_fn_inbuilt = clf.decision_function(X_cv)

```

561

In [87]:

```

print("Function based decision function values are",dec_fn_imlem[:5])
print("sklearn based decision function values are ",dec_fn_inbuilt.tolist()[:5])

```

```

Function based decision function values are [-4.161095825122163, 1.7233399647584635, -
2.6499005585185555, 1.3091961267094492, -0.3664902922009814]
sklearn based decision function values are [-4.161095825122163, 1.7233399647584635, -
2.6499005585185555, 1.3091961267094492, -0.3664902922009814]

```

we can observe that both implementations result in the same values of decision function.

From above we can observe that the function implementation result and inbuilt scikitlearn function results are same.

8F: Implementing Platt Scaling to find $P(Y=1|X)$

Check this [PDF](#)

Let the output of a learning method be $f(x)$. To get calibrated probabilities, pass the output through a sigmoid:

$$P(y = 1|f) = \frac{1}{1 + \exp(Af + B)} \quad (1)$$

where the parameters A and B are fitted using maximum likelihood estimation from a fitting training set (f_i, y_i) . Gradient descent is used to find A and B such that they are the solution to:

$$\underset{A,B}{\operatorname{argmin}} \left\{ - \sum_i y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \right\}, \quad (2)$$

where

$$p_i = \frac{1}{1 + \exp(Af_i + B)} \quad (3)$$

Two questions arise: where does the sigmoid train set come from? and how to avoid overfitting to this training set?

If we use the same data set that was used to train the model we want to calibrate, we introduce unwanted bias. For example, if the model learns to discriminate the train set perfectly and orders all the negative examples before the positive examples, then the sigmoid transformation will output just a 0,1 function. So we need to use an independent calibration set in order to get good posterior probabilities. This, however, is not a draw back, since the same set can be used for model and parameter selection.

To avoid overfitting to the sigmoid train set, an out-of-sample model is used. If there are N_+ positive examples and N_- negative examples in the train set, for each training example Platt Calibration uses target values u_+ and u_-

(instead of 1 and 0, respectively), where

$$y_+ = \frac{N_+ + 1}{N_+ + 2}; y_- = \frac{1}{N_- + 2} \quad (4)$$

For a more detailed treatment, and a justification of these particular target values see (Platt, 1999).

TASK F

1. Apply SGD algorithm with (f_{CV}, y_{CV}) and find the weight W intercept b Note: here our data is of one dimensional so we will have a one dimensional weight vector i.e $W.shape (1,)$

Note1: Don't forget to change the values of y_{CV} as mentioned in the above image. you will calculate y_+, y_- based on data points in train data

Note2: the Sklearn's SGD algorithm doesn't support the real valued outputs, you need to use the code that was done in the 'Logistic Regression with SGD and L2' Assignment after modifying loss function, and use same parameters that used in that assignment.

```
def log_loss(w, b, X, Y):
    N = len(X)
    sum_log = 0
    for i in range(N):
        sum_log += Y[i]*np.log10(sig(w, X[i], b)) + (1-Y[i])*np.log10(1-sig(w, X[i], b))
    return -1*sum_log/N
```

if $Y[i]$ is 1, it will be replaced with y_+ value else it will be replaced with y_- value

1. For a given data point from X_{test} , $P(Y = 1 | X) = \frac{1}{1 + \exp(-(W * f_{test} + b))}$ where $f_{test} = \text{decision_function}(X_{test})$, W and b will be learned as mentioned in the above step

In [88]:

```
def logloss(y_true, y_pred):
    '''In this function, we will compute log loss '''

    n = len(y_true)
    x = np.log10(y_pred)
    x1 = np.log10(np.ones_like(y_pred) - y_pred)
    loss = 0
    for j in range(n):
        loss = loss + ((y_true[j]*x[j]) + ((1-y_true[j])*x1[j]))
    loss*=(-1/n)
    return loss
```

In [89]:

```
def initialize_weights(dim):
    ''' In this function, we will initialize our weights and bias'''
    #initialize the weights to zeros array of (1,dim) dimensions
    #you use zeros_like function to initialize zero, check this link https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
    #initialize bias to zero

    w = np.zeros_like(dim)
    b = 0

    return w,b
```

In [90]:

```
def sigmoid(z):
    ''' In this function, we will return sigmoid of z'''
    # compute sigmoid(z) and return

    sigma = (1/(1+np.exp(-z)))

    return sigma
```

In [91]:

```
def gradient_dw(x,y,w,b,alpha,N):
    '''In this function, we will compute the gradient w.r.to w'''

    dw = (x*(y - sigmoid(np.matmul(w,x) + b))) - ((alpha/N)*w)

    return dw
```

In [92]:

```
def gradient_db(x,y,w,b):
    '''In this function, we will compute gradient w.r.to b'''

    db = (y - sigmoid(np.matmul(w,x) + b))

    return db
```

In [93]:

```
def train(X_train,y_train,epochs,alpha,eta0):
    ''' In this function, we will implement logistic regression'''
    #Here eta0 is learning rate
    #implement the code as follows
    # initialize the weights (call the initialize_weights(X_train[0]) function)
    # for every epoch
        # for every data point(X_train,y_train)
            #compute gradient w.r.to w (call the gradient_dw() function)
            #compute gradient w.r.to b (call the gradient_db() function)
            #update w, b
        # predict the output of x_train[for all data points in X_train] using w,b
        #compute the loss between predicted and actual values (call the loss function
    )

        # store all the train loss values in a list
        # predict the output of x_test[for all data points in X_test] using w,b
        #compute the loss between predicted and actual values (call the loss function
    )

        # store all the test loss values in a list
        # you can also compare previous loss and current loss, if loss is not updating then stop the process and return w,b

    w,b = initialize_weights(X_train[0])
    N = len(X_train)

    loss_train = []
    for i in range(epochs):
        w_prev,b_prev = w,b
        y_pred_train = []
        for j in range(len(X_train)):
            dw = gradient_dw(X_train[j,:], y_train[j], w, b, alpha, N)
            db = gradient_db(X_train[j,:], y_train[j], w, b)

            w = w + (eta0 * dw)
            b = b + (eta0 * db)
        for j in range(len(X_train)):
            y_pred_train.append(sigmoid(np.matmul(w,X_train[j,:]) + b))

    present=logloss(y_train, y_pred_train)
    loss_train.append(present)
```

```

#print(present)

if(i!=0 and loss_train[i-1]<present):
    print("Minimum loss achieved is",loss_train[i-1])
    return w_prev,b_prev,loss_train[:i]
print("Minimum loss achieved is",present)
return w,b,loss_train

```

In [94]:

```

N_pos = np.count_nonzero(Y_tr)
N_neg = len(Y_tr)-N_pos
y_pos = (N_pos+1)/(N_pos+2)
y_neg = (1/(N_neg+2))
#print(N_pos,N_neg,y_pos,y_neg)
print(Y_cv.tolist()[0:4])
Y_cv_mod = np.zeros_like(Y_cv,dtype =float)
for i in range(len(Y_cv)):
    if Y_cv[i]==1:
        Y_cv_mod[i] = y_pos
    else:
        Y_cv_mod[i] = y_neg
print(Y_cv_mod.tolist()[0:4])

```

```

[0, 1, 0, 1]
[0.00048192771084337347, 0.9989235737351991, 0.00048192771084337347, 0.9989235737351991]

```

In the above results, list-1 indicates true or actual labels of Y_cv. list-2 indicates the modified labels of Y_cv according to the given formulae.

We can observe that 0 label is replaced with 0.0004766 and label 1 is replaced with 0.9988962

In [95]:

```

alpha=0.0001
eta0=0.001
N=len(X_tr)
epochs=1000
#print(np.array(dec_fn_imlem).reshape(-1,1),np.array(Y_cv_mod).reshape(-1,1))
f_cv = np.array(dec_fn_imlem).reshape(-1,1)
Y_cv_mod = np.array(Y_cv_mod).reshape(-1,1)
w,b,train_loss = train(f_cv,Y_cv_mod,epochs,alpha,eta0)

```

Minimum loss achieved is [0.07559336]

In [96]:

```


import matplotlib.pyplot as plt


epochs=[i for i in range(len(train_loss))]
plt.plot(epochs,train_loss,label= 'Train Loss')

plt.scatter(epochs,train_loss)

plt.legend()
plt.title('Epochs vs Logistic Loss graph')
plt.xlabel('epochs')
plt.ylabel('Logistic loss')
plt.show()

```





In [97]:

```
f_test = decision_function(X_test,sv,y_alpha[0],intercept,gam)
print(f_test[:4])
proba_1=[]
for i in range(len(X_test)):
    proba_1.append(sigmoid(np.dot(w,f_test[i]) + b))

proba_1 = [i.tolist() for i in proba_1]
proba_1 = [i[0] for i in proba_1]
print(proba_1[:4])
```

```
[-2.392145229826883, -2.9514350778685126, -1.6627732024593085, -4.619833147669826]
[0.011427693208354434, 0.003914680589311406, 0.04507717531280103, 0.000157258920370584
34]
```

we can observe that if decision function value is negative it's probability of being class 1 $P(y = 1/x) < 0.5$ and is low and if it is positive it's $P(y = 1/x) > 0.5$ and high.

Note: in the above algorithm, the steps 2, 4 might need hyper parameter tuning, To reduce the complexity of the assignment we are excluding the hyperparameter tuning part, but interested students can try that

If any one wants to try other calibration algorithm isotonic regression also please check these tutorials

1. <http://fa.bianp.net/blog/tag/scikit-learn.html#fn:1>
2. https://drive.google.com/open?id=1MzmA7QaP58RDzocBORBmRiWfI7Co_VJ7
3. https://drive.google.com/open?id=133odBinMOIVb_rh_GQxxsyMRyW-Zts7a
4. https://stat.fandom.com/wiki/Isotonic_regression#Pool_Adjacent_Violators_Algorithm